

Operating-System Structures





Contents

- n What Operating Systems Do
- n Computer-System Organization
- n Computer-System Architecture
- n Operating-System Structure
- n Operating-System Operations
- n Operating System Services/ Functions
- n User Operating System Interface
- n System Calls
- n Types of System Calls
- n System Programs
- n System Boot





What is an Operating System?

- n A program that acts as an intermediary between a user of a computer and the computer hardware
- n Operating system goals:
 - | Execute user programs and make solving user problems easier
 - | Make the computer system convenient to use
 - | Use the computer hardware in an efficient manner





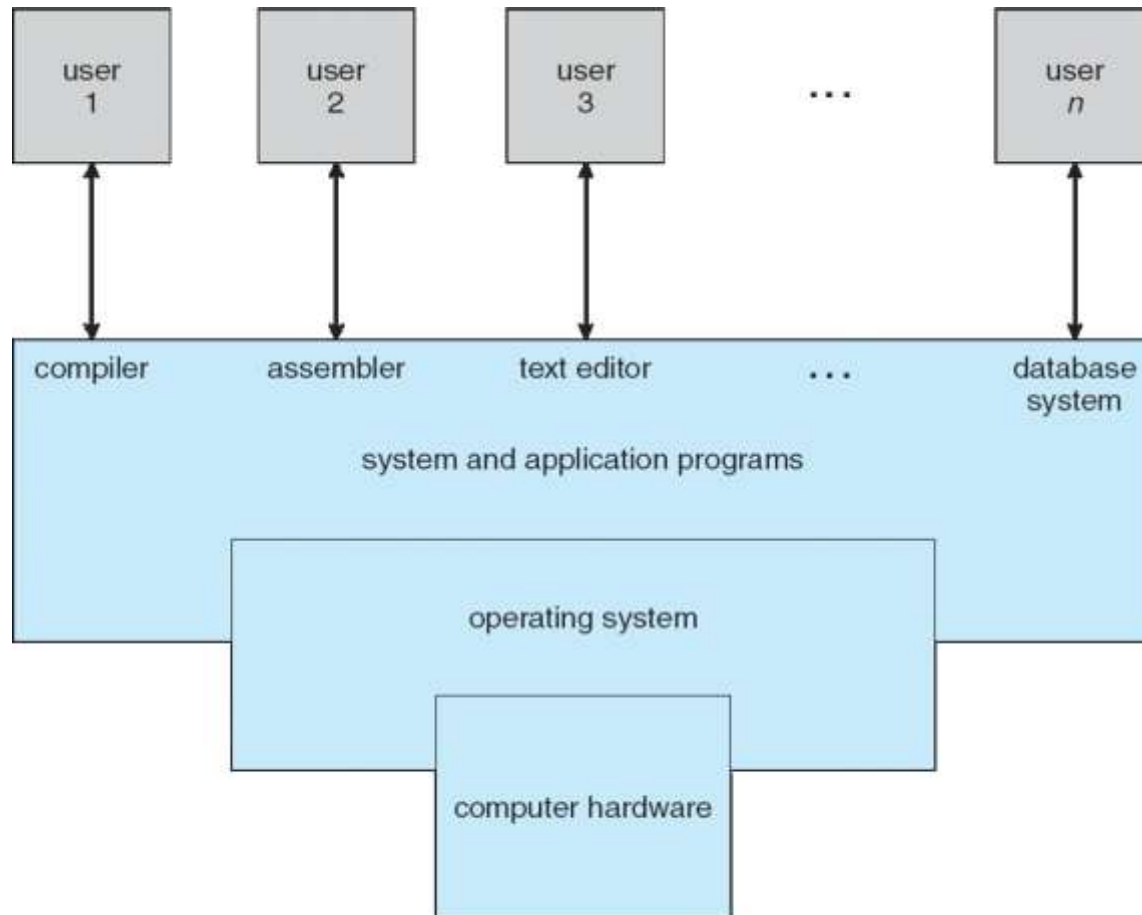
Computer System Structure

- n Computer system can be divided into four components:
 - | Hardware – provides basic computing resources
 - ▶ CPU, memory, I/O devices
 - | Operating system
 - ▶ Controls and coordinates use of hardware among various applications and users
 - | Application programs – define the ways in which the system resources are used to solve the computing problems of the users
 - ▶ Word processors, compilers, web browsers, database systems, video games
 - | Users
 - ▶ People, machines, other computers





Four Components of a Computer System





What Operating Systems Do

- n Depends on the point of view
- n Users want convenience, **ease of use** and **good performance**
 - l Don't care about **resource utilization**
- n But shared computer such as **mainframe** or **minicomputer** must keep all users happy
- n Users of dedicate systems such as **workstations** have dedicated resources but frequently use shared resources from **servers**
- n Handheld computers are resource poor, optimized for usability and battery life
- n Some computers have little or no user interface, such as embedded computers in devices and automobiles





Operating System Definition

- n OS is a **resource allocator**
 - | Manages all resources
 - | Decides between conflicting requests for efficient and fair resource use
- n OS is a **control program**
 - | Controls execution of programs to prevent errors and improper use of the computer





Operating System Definition (Cont.)

- n No universally accepted definition
- n “Everything a vendor ships when you order an operating system” is a good approximation
 - | But varies wildly
- n “The one program running at all times on the computer” is the **kernel**.
- n Everything else is either
 - | a system program (ships with the operating system) , or
 - | an application program.





Computer Startup


- n **bootstrap program** is loaded at power-up or reboot
 - | Typically stored in ROM or EPROM, generally known as **firmware**
 - | Initializes all aspects of system
 - | Loads operating system kernel and starts execution





Computer System Organization

- n Computer-system operation
 - | One or more CPUs, device controllers connect through common bus providing access to shared memory
 - | Concurrent execution of CPUs and devices competing for memory cycles





Computer-System Operation

- n I/O devices and the CPU can execute concurrently
- n Each device controller is in charge of a particular device type
- n Each device controller has a local buffer
- n CPU moves data from/to main memory to/from local buffers
- n I/O is from the device to local buffer of controller
- n Device controller informs CPU that it has finished its operation by causing an **interrupt**





Common Functions of Interrupts

- n Interrupt transfers control to the interrupt service routine generally, through the **interrupt vector**, which contains the addresses of all the service routines
- n Interrupt architecture must save the address of the interrupted instruction
- n A **trap** or **exception** is a software-generated interrupt caused either by an error or a user request
- n An operating system is **interrupt driven**





Operating-System Operations

- n **Interrupt driven** (hardware and software)
 - | Hardware interrupt by one of the devices
 - | Software interrupt (**exception** or **trap**):
 - ▶ Software error (e.g., division by zero)
 - ▶ Request for operating system service
 - ▶ Other process problems include infinite loop, processes modifying each other or the operating system





Operating-System Operations (cont.)

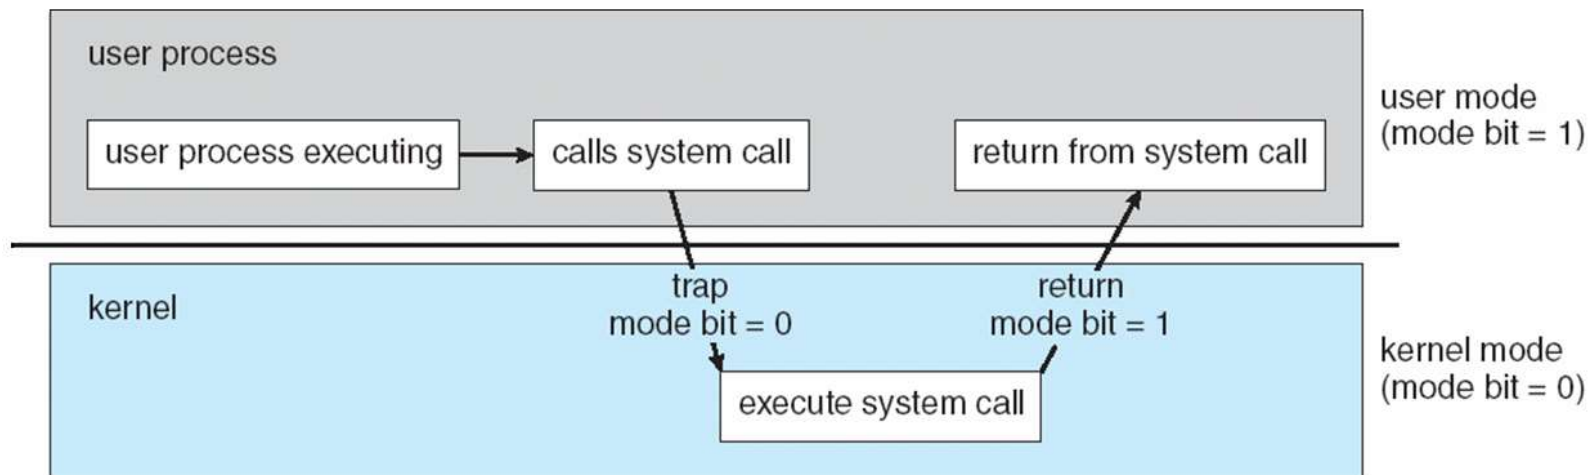
- n **Dual-mode** operation allows OS to protect itself and other system components
 - | **User mode** and **kernel mode**
 - | **Mode bit** provided by hardware
 - ▶ Provides ability to distinguish when system is running user code or kernel code
 - ▶ Some instructions designated as **privileged**, only executable in kernel mode
 - ▶ System call changes mode to kernel, return from call resets it to user
- n Increasingly CPUs support multi-mode operations
 - | i.e. **virtual machine manager (VMM)** mode for guest **VMs**





Transition from User to Kernel Mode

- n Timer to prevent infinite loop / process hogging resources
 - | Timer is set to interrupt the computer after some time period
 - | Keep a counter that is decremented by the physical clock.
 - | Operating system set the counter (privileged instruction)
 - | When counter zero generate an interrupt
 - | Set up before scheduling process to regain control or terminate program that exceeds allotted time





Operating System Services

- n Operating systems provide an environment for execution of programs and services to programs and users
- n One set of operating-system services provides functions that are helpful to the user:
 - | **User interface** - Almost all operating systems have a user interface (**UI**).
 - ▶ Varies between **Command-Line (CLI)**, **Graphics User Interface (GUI)**, **Batch**
 - | **Program execution** - The system must be able to load a program into memory and to run that program, end execution, either normally or abnormally (indicating error)
 - | **I/O operations** - A running program may require I/O, which may involve a file or an I/O device





Operating System Services (Cont.)

- n One set of operating-system services provides functions that are helpful to the user (Cont.):
 - | **File-system manipulation** - The file system is of particular interest. Programs need to read and write files and directories, create and delete them, search them, list file information, permission management.
 - | **Communications** – Processes may exchange information, on the same computer or between computers over a network
 - ▶ Communications may be via shared memory or through message passing (packets moved by the OS)
 - | **Error detection** – OS needs to be constantly aware of possible errors
 - ▶ May occur in the CPU and memory hardware, in I/O devices, in user program
 - ▶ For each type of error, OS should take the appropriate action to ensure correct and consistent computing
 - ▶ Debugging facilities can greatly enhance the user's and programmer's abilities to efficiently use the system





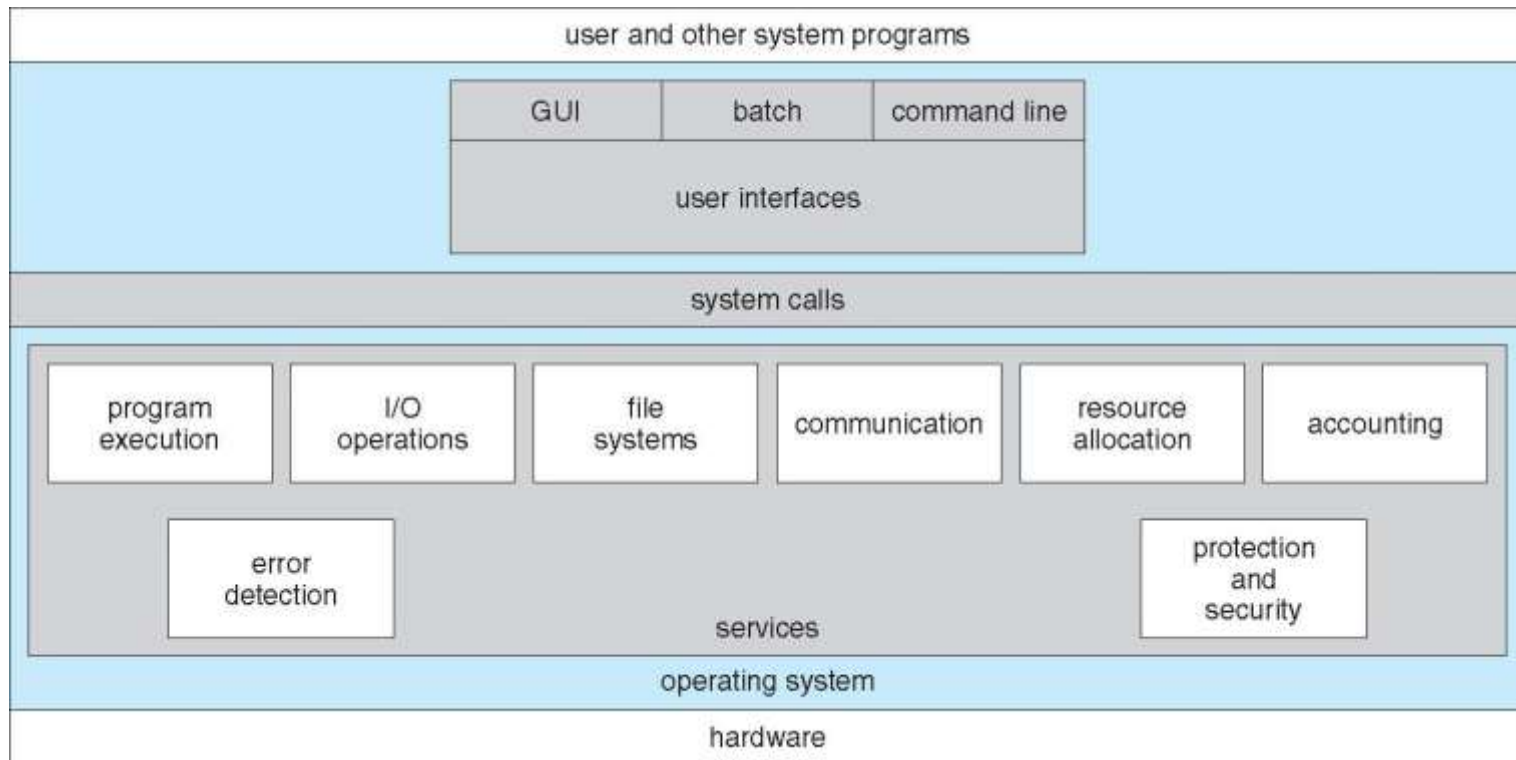
Operating System Services (Cont.)

- n Another set of OS functions exists for ensuring the efficient operation of the system itself via resource sharing
 - | **Resource allocation** - When multiple users or multiple jobs running concurrently, resources must be allocated to each of them
 - ▶ Many types of resources - CPU cycles, main memory, file storage, I/O devices.
 - | **Accounting** - To keep track of which users use how much and what kinds of computer resources
 - | **Protection and security** - The owners of information stored in a multiuser or networked computer system may want to control use of that information, concurrent processes should not interfere with each other
 - ▶ **Protection** involves ensuring that all access to system resources is controlled
 - ▶ **Security** of the system from outsiders requires user authentication, extends to defending external I/O devices from invalid access attempts




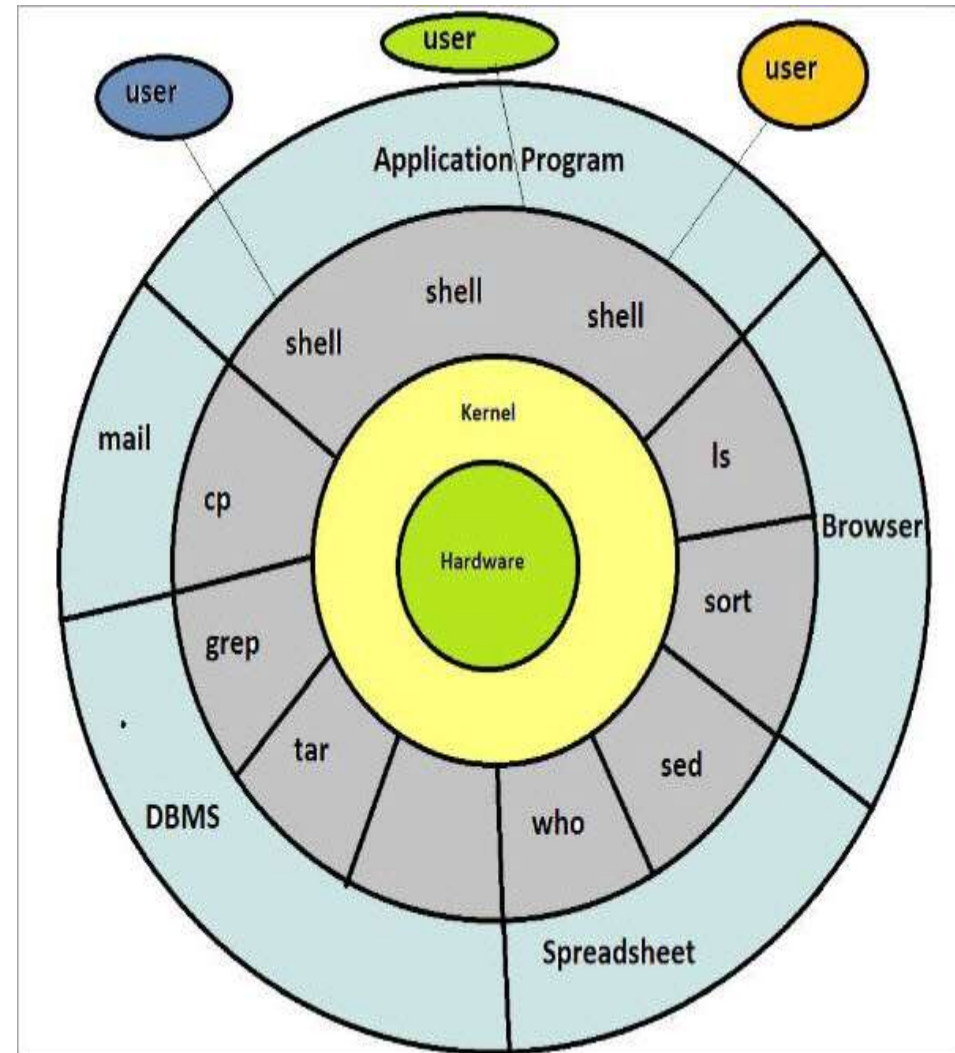



A View of Operating System Services





Kernel in windows and linux





User Operating System Interface - CLI

CLI or **command interpreter** allows direct command entry

- | Sometimes implemented in kernel, sometimes by systems program
- | Sometimes multiple flavors implemented – **shells**
- | Primarily fetches a command from user and executes it
- | Sometimes commands built-in, sometimes just names of programs
 - ▶ If the latter, adding new features doesn't require shell modification





Bourne Shell Command Interpreter

```
Default
New Info Close Execute Bookmarks

PBGMac-Pro:~ pbg$ w
15:24 up 56 mins, 2 users, load averages: 1.51 1.53 1.65
USER      TTY      FROM          LOGIN@  IDLE   WHAT
pbg       console -              14:34   50    -
pbg       s000    -              15:05   -    w
PBGMac-Pro:~ pbg$ iostat 5
            disk0      disk1      disk10      cpu      load average
      KB/t tps  MB/s    KB/t tps  MB/s    KB/t tps  MB/s  us sy id  1m  5m  15m
    33.75 343 11.30    64.31 14  0.88    39.67  0  0.02  11  5 84  1.51 1.53 1.65
     5.27 320  1.65     0.00  0  0.00     0.00  0  0.00   4  2 94  1.39 1.51 1.65
     4.28 329  1.37     0.00  0  0.00     0.00  0  0.00   5  3 92  1.44 1.51 1.65
^C
PBGMac-Pro:~ pbg$ ls
Applications          Music                  WebEx
Applications (Parallels)  Pando Packages       config.log
Desktop               Pictures              getsmartdata.txt
Documents             Public                imp
Downloads             Sites                 log
Dropbox               Thumbs.db             panda-dist
Library              Virtual Machines     prob.txt
Movies               Volumes               scripts
PBGMac-Pro:~ pbg$ pwd
/Users/pbg
PBGMac-Pro:~ pbg$ ping 192.168.1.1
PING 192.168.1.1 (192.168.1.1): 56 data bytes
64 bytes from 192.168.1.1: icmp_seq=0 ttl=64 time=2.257 ms
64 bytes from 192.168.1.1: icmp_seq=1 ttl=64 time=1.262 ms
^C
--- 192.168.1.1 ping statistics ---
2 packets transmitted, 2 packets received, 0.0% packet loss
round-trip min/avg/max/stddev = 1.262/1.760/2.257/0.498 ms
PBGMac-Pro:~ pbg$
```





User Operating System Interface - GUI


- n User-friendly **desktop** metaphor interface
 - | Usually mouse, keyboard, and monitor
 - | **Icons** represent files, programs, actions, etc
 - | Various mouse buttons over objects in the interface cause various actions (provide information, options, execute function, open directory (known as a **folder**))
 - | Invented at Xerox PARC
- n Many systems now include both CLI and GUI interfaces
 - | Microsoft Windows is GUI with CLI “command” shell
 - | Apple Mac OS X is “Aqua” GUI interface with UNIX kernel underneath and shells available
 - | Unix and Linux have CLI with optional GUI interfaces (CDE, KDE, GNOME)






Touchscreen Interfaces

- n Touchscreen devices require new interfaces
 - | Mouse not possible or not desired
 - | Actions and selection based on gestures
 - | Virtual keyboard for text entry
- | Voice commands.





The Mac OS X GUI





System Calls

- n Programming interface to the services provided by the OS
- n Typically written in a high-level language (C or C++)
- n Mostly accessed by programs via a high-level **Application Programming Interface (API)** rather than direct system call use
- n Three most common APIs are Win32 API for Windows, POSIX API for POSIX-based systems (including virtually all versions of UNIX, Linux, and Mac OS X), and Java API for the Java virtual machine (JVM)

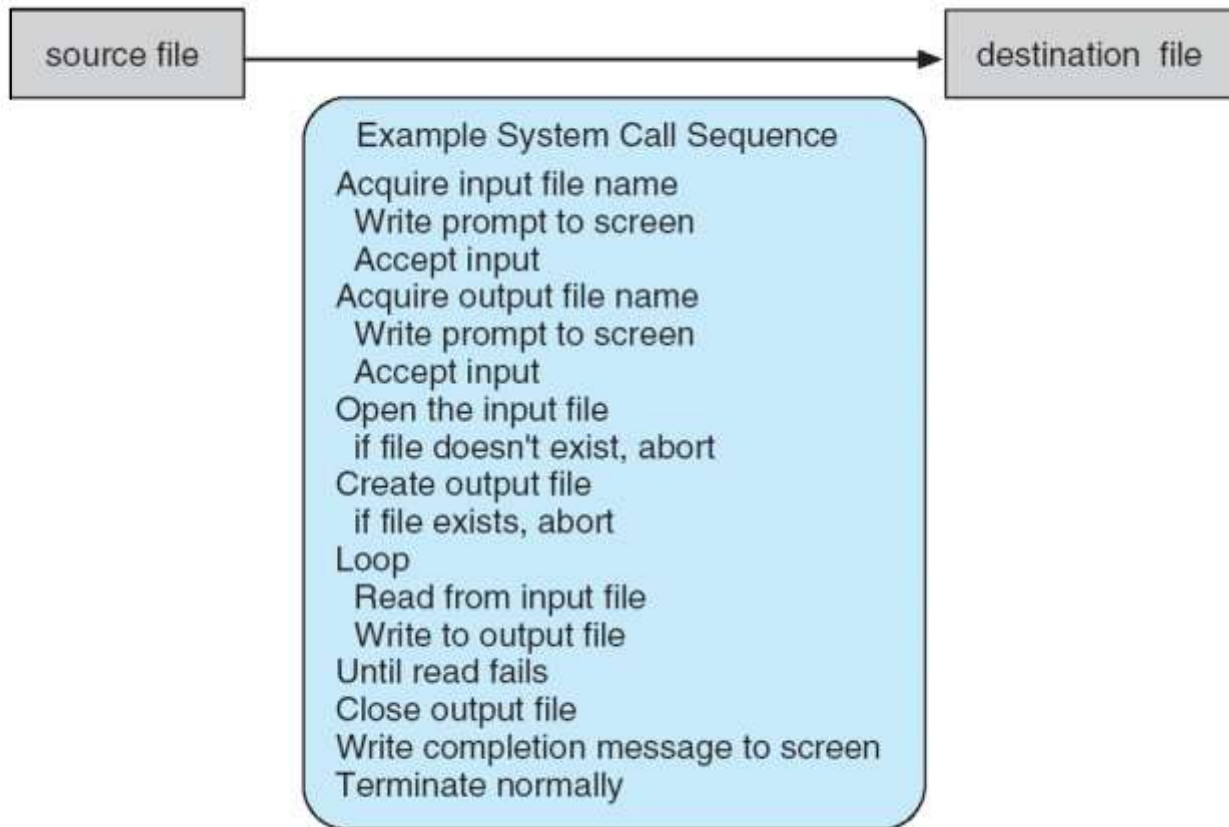
Note that the system-call names used throughout this text are generic





Example of System Calls

- n System call sequence to copy the contents of one file to another file





Example of Standard API

EXAMPLE OF STANDARD API

As an example of a standard API, consider the `read()` function that is available in UNIX and Linux systems. The API for this function is obtained from the `man` page by invoking the command

```
man read
```

on the command line. A description of this API appears below:

```
#include <unistd.h>

ssize_t  read(int fd, void *buf, size_t count)
```

return	function	parameters
value	name	

A program that uses the `read()` function must include the `unistd.h` header file, as this file defines the `ssize_t` and `size_t` data types (among other things). The parameters passed to `read()` are as follows:

- `int fd`—the file descriptor to be read
- `void *buf`—a buffer where the data will be read into
- `size_t count`—the maximum number of bytes to be read into the buffer

On a successful read, the number of bytes read is returned. A return value of 0 indicates end of file. If an error occurs, `read()` returns `-1`.






System Call Implementation

- n Typically, a number associated with each system call
 - | **System-call interface** maintains a table indexed according to these numbers
- n The system call interface invokes the intended system call in OS kernel and returns status of the system call and any return values
- n The caller need know nothing about how the system call is implemented
 - | Just needs to obey API and understand what OS will do as a result call
 - | Most details of OS interface hidden from programmer by API
 - ▶ Managed by run-time support library (set of functions built into libraries included with compiler)





API – System Call – OS Relationship





System Call Parameter Passing

- n Often, more information is required than simply identity of desired system call
 - | Exact type and amount of information vary according to OS and call
- n Three general methods used to pass parameters to the OS
 - | Simplest: pass the parameters in registers
 - ▶ In some cases, may be more parameters than registers
 - | Parameters stored in a block, or table, in memory, and address of block passed as a parameter in a register
 - ▶ This approach taken by Linux and Solaris
 - | Parameters placed, or **pushed**, onto the **stack** by the program and **popped** off the stack by the operating system
 - | Block and stack methods do not limit the number or length of parameters being passed





Types of System Calls

- n Process control
 - | create process, terminate process
 - | end, abort
 - | load, execute
 - | get process attributes, set process attributes
 - | wait for time
 - | wait event, signal event
 - | allocate and free memory
 - | Dump memory if error
 - | **Debugger** for determining **bugs, single step** execution
 - | **Locks** for managing access to shared data between processes





Types of System Calls

- n File management
 - | create file, delete file
 - | open, close file
 - | read, write, reposition
 - | get and set file attributes
- n Device management
 - | request device, release device
 - | read, write, reposition
 - | get device attributes, set device attributes
 - | logically attach or detach devices





Types of System Calls (Cont.)

- n Information maintenance
 - | get time or date, set time or date
 - | get system data, set system data
 - | get and set process, file, or device attributes
- n Communications
 - | create, delete communication connection
 - | send, receive messages if **message passing model** to **host name** or **process name**
 - ▶ From **client** to **server**
 - | **Shared-memory model** create and gain access to memory regions
 - | transfer status information
 - | attach and detach remote devices





Types of System Calls (Cont.)

- n Protection
 - | Control access to resources
 - | Get and set permissions
 - | Allow and deny user access





Examples of Windows and Unix System Calls


	Windows	Unix
Process Control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File Manipulation	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device Manipulation	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information Maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communication	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shmget() mmap()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()





Standard C Library Example


- n C program invoking printf() library call, which calls write() system call





Layered Approach

- n The operating system is divided into a number of layers (levels), each built on top of lower layers. The bottom layer (layer 0), is the hardware; the highest (layer N) is the user interface.
- n With modularity, layers are selected such that each uses functions (operations) and services of only lower-level layers





System Boot

- n When power initialized on system, execution starts at a fixed memory location
 - | Firmware ROM used to hold initial boot code
- n Operating system must be made available to hardware so hardware can start it
 - | Small piece of code – **bootstrap loader**, stored in **ROM** or **EEPROM** locates the kernel, loads it into memory, and starts it
 - | Sometimes two-step process where **boot block** at fixed location loaded by ROM code, which loads bootstrap loader from disk
- n Common bootstrap loader, **GRUB**, allows selection of kernel from multiple disks, versions, kernel options
- n Kernel loads and system is then **running**

